

# Computer Organization and Design

A decorative graphic consisting of a vertical blue line on the left side of the slide, a horizontal blue line intersecting it, and a horizontal light blue gradient bar extending to the right from the intersection point.


## Arithmetic for Computers

# Review: MIPS (RISC) Design Principles


- **Simplicity favors regularity**
  - fixed size instructions
  - small number of instruction formats
  - opcode always the first 6 bits
- **Smaller is faster**
  - limited instruction set
  - limited number of registers in register file
  - limited number of addressing modes
- **Make the common case fast**
  - arithmetic operands from the register file (load-store machine)
  - allow instructions to contain immediate operands
- **Good design demands good compromises**
  - three instruction formats

# Review: MIPS Addressing Modes


## 1. Register addressing




## 2. Base (displacement) addressing




## 3. Immediate addressing



## 4. PC-relative addressing



## 5. Pseudo-direct addressing



# Review: Number Representations

## □ 32-bit signed numbers (2's complement):

0000	0000	0000	0000	0000	0000	0000	0000	0	$= 0_{\text{ten}}$	
0000	0000	0000	0000	0000	0000	0000	0000	1	$= + 1_{\text{ten}}$	<i>maxint</i>
...										
0111	1111	1111	1111	1111	1111	1111	1111	0	$= + 2,147,483,646_{\text{ten}}$	
0111	1111	1111	1111	1111	1111	1111	1111	1	$= + 2,147,483,647_{\text{ten}}$	
1000	0000	0000	0000	0000	0000	0000	0000	0	$= - 2,147,483,648_{\text{ten}}$	<i>minint</i>
1000	0000	0000	0000	0000	0000	0000	0000	1	$= - 2,147,483,647_{\text{ten}}$	
...										
1111	1111	1111	1111	1111	1111	1111	1111	0	$= - 2_{\text{ten}}$	
1111	1111	1111	1111	1111	1111	1111	1111	1	$= - 1_{\text{ten}}$	

MSB (Most Significant Bit) is indicated by a red oval around the first bit of each row.

LSB (Least Significant Bit) is indicated by a red oval around the last bit of each row.

## □ Converting <32-bit values into 32-bit values

- 1 copy the most significant bit (the sign bit) into the “empty” bits

0010 → 0000 0010

1010 → 1111 1010

- 1 **sign extend** versus zero extend (lb vs. lbu)

# Arithmetic for Computers

- Operations on integers
  - Addition and subtraction
  - Multiplication and division
- What about fractions and real numbers?
  - Representation and operations
- How are overflow scenarios handled?
  - e.g. An operation creates a number bigger than can be represented
- How does hardware really multiply and divide numbers?

# MIPS Arithmetic Logic Unit (ALU)

- Must support the Arithmetic/Logic operations of the ISA

`add, addi, addiu, addu`

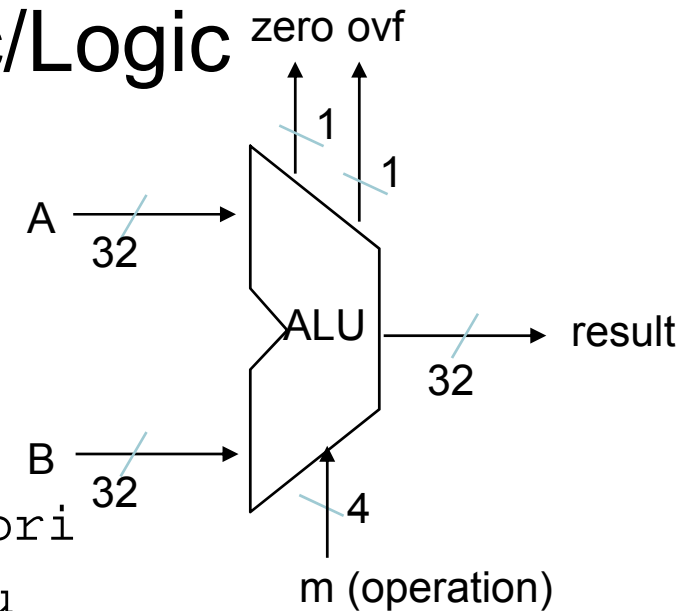
`sub, subu`

`mult, multu, div, divu`

`sqr`

`and, andi, nor, or, ori, xor, xori`

`beq, bne, slt, slti, sltiu, sltu`



- With special handling for

**sign extend** – `addi, addiu, slti, sltiu`


**zero extend** – `andi, ori, xori`

**overflow detection** – `add, addi, sub`

**Review Appendix C (from CD or lecture page) for more details on ALU design**

# Integer Addition

## ■ Example: $7 + 6$



## ■ Overflow if result out of range

- Adding +ve and -ve operands, no overflow
- Adding two +ve operands
  - Overflow if result sign bit is 1
- Adding two -ve operands
  - Overflow if result sign bit is 0

# Integer Subtraction

- Add negation of second operand

- Example:  $7 - 6 = 7 + (-6)$

+7:	0000 0000 ... 0000 0111
-6:	1111 1111 ... 1111 1010
<hr/>	
+1:	0000 0000 ... 0000 0001

- **Overflow if result out of range**

- Subtracting two +ve or two -ve operands, no overflow
- Subtracting +ve from -ve operand
  - Overflow if result sign bit is 0
- Subtracting -ve from +ve operand
  - Overflow if result sign bit is 1



# Dealing with Overflow

- Some languages (e.g., C) ignore overflow
  - C compilers use MIPS **addu**, **addui**, **subu** instructions
- Other languages (e.g., Ada, Fortran) require raising an exception
  - Use MIPS **add**, **addi**, **sub** instructions
  - On overflow, invoke exception handler
    - Save PC in exception program counter (EPC) register
    - Jump to predefined handler address
    - **mfhc0** (move from coprocessor reg) instruction can retrieve EPC value, to return after corrective action

# Arithmetic for Multimedia

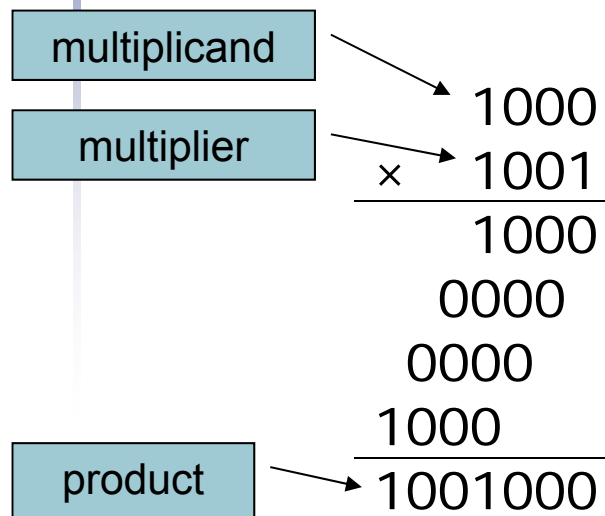
- Graphics and media processing operates on vectors of 8-bit and 16-bit data
  - Use 64-bit adder, with partitioned carry chain
    - Operate on 8×8-bit, 4×16-bit, or 2×32-bit vectors
  - SIMD (single-instruction, multiple-data)
- Saturating operations
  - On overflow, result is set to the largest representable value
    - c.f. 2s-complement modulo arithmetic
  - E.g., clipping in audio, saturation in video

$$\begin{array}{r} 10000000 \\ + 10000000 \\ \hline 100000000 \end{array}$$

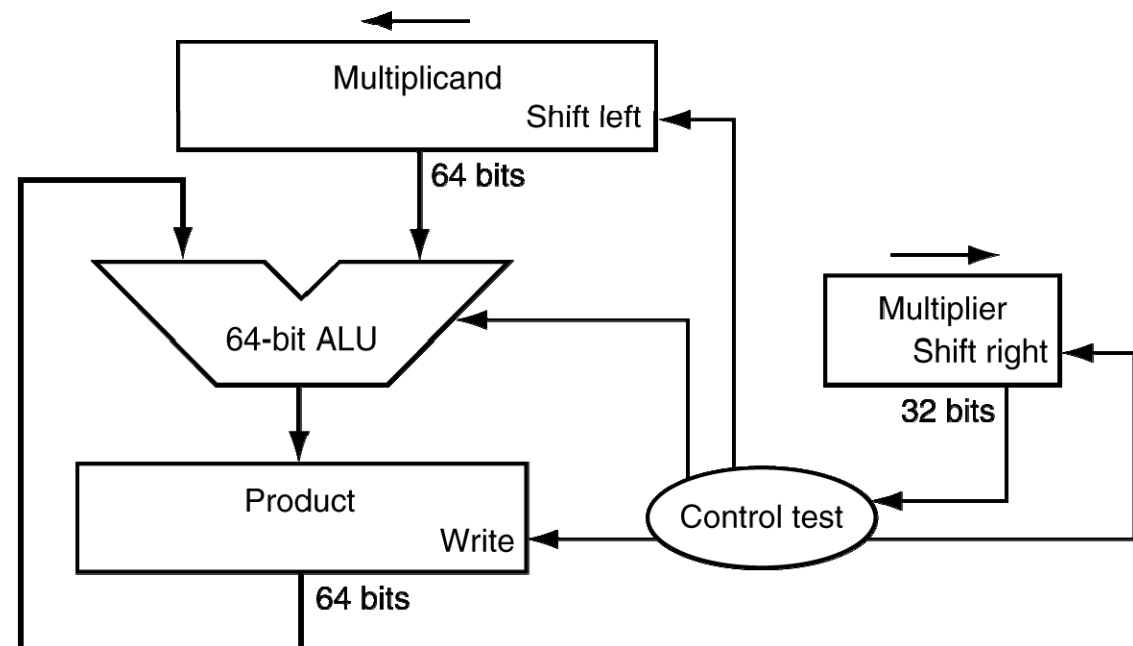
11111111

# Multiplication

- Start with long-multiplication approach



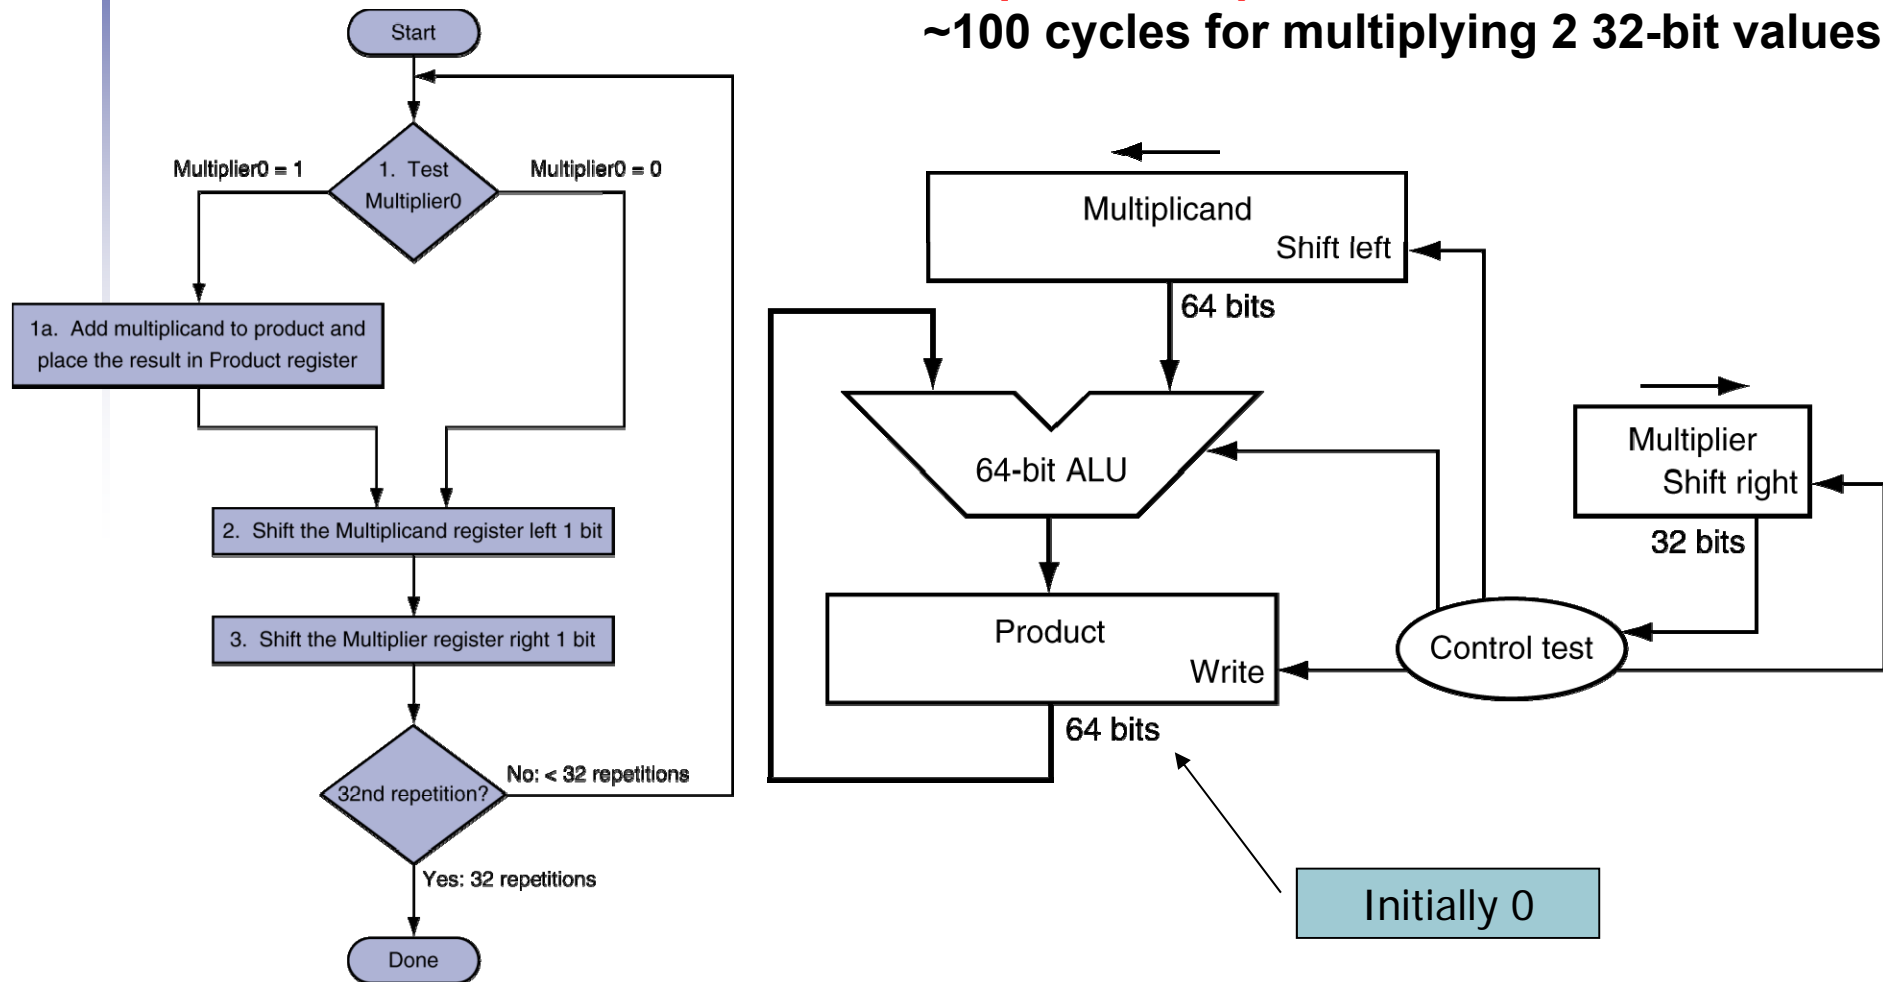
Length of product is the sum of operand lengths



# Multiplication Hardware

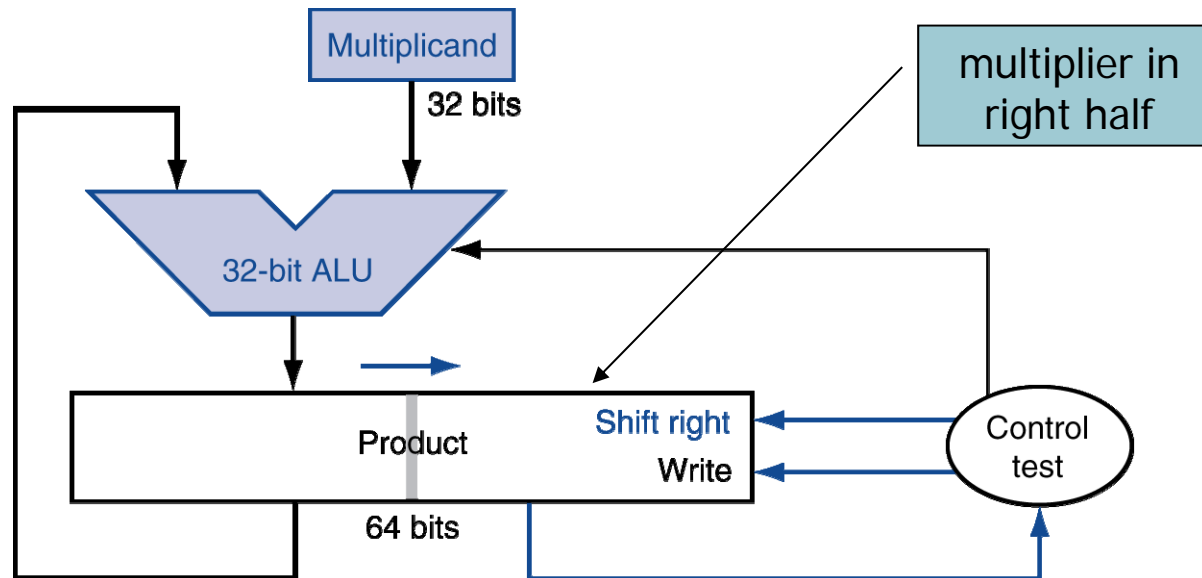
## Sequential Implementation

~100 cycles for multiplying 2 32-bit values



# Optimized Multiplier


- Perform steps in parallel: add/shift



- One cycle per partial-product addition
  - That's ok, if frequency of multiplications is low
  - ~32 cycles for multiplying 2 32-bit values

# Even Faster Multiplier

- Uses multiple adders
  - Cost/performance tradeoff
    - 31 adders  $\rightarrow \log_2(31) = \sim 5$  cycles for multiplying 2 32-bit values

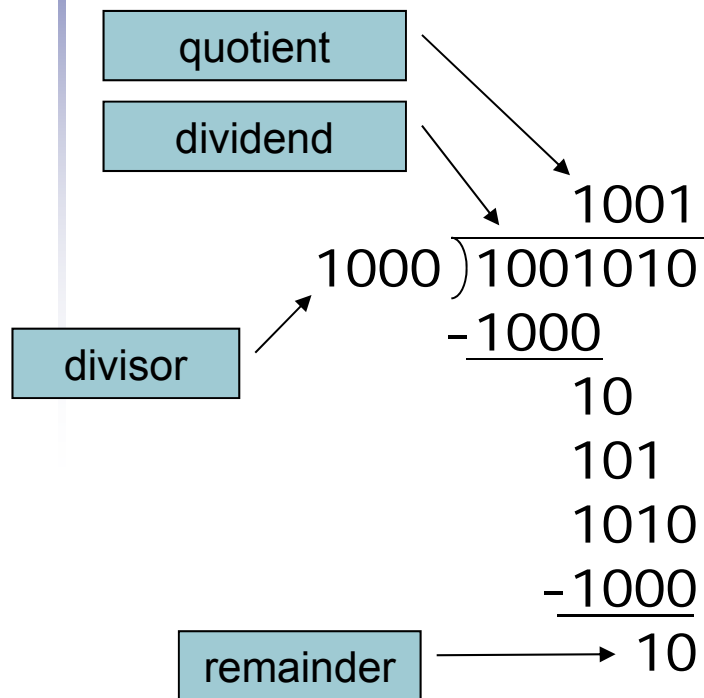


- Can be pipelined
  - Several multiplication performed in parallel

# MIPS Multiplication

- Two 32-bit registers for product
  - **HI**: most-significant 32 bits
  - **LO**: least-significant 32-bits
- Instructions
  - `mult rs, rt` / `multu rs, rt`
    - 64-bit product in HI/LO
  - `mfhi rd` / `mflo rd`
    - Move from HI/LO to rd
    - Can test HI value to see if product overflows 32 bits
  - `mul rd, rs, rt`
    - Least-significant 32 bits of product → rd

# Division

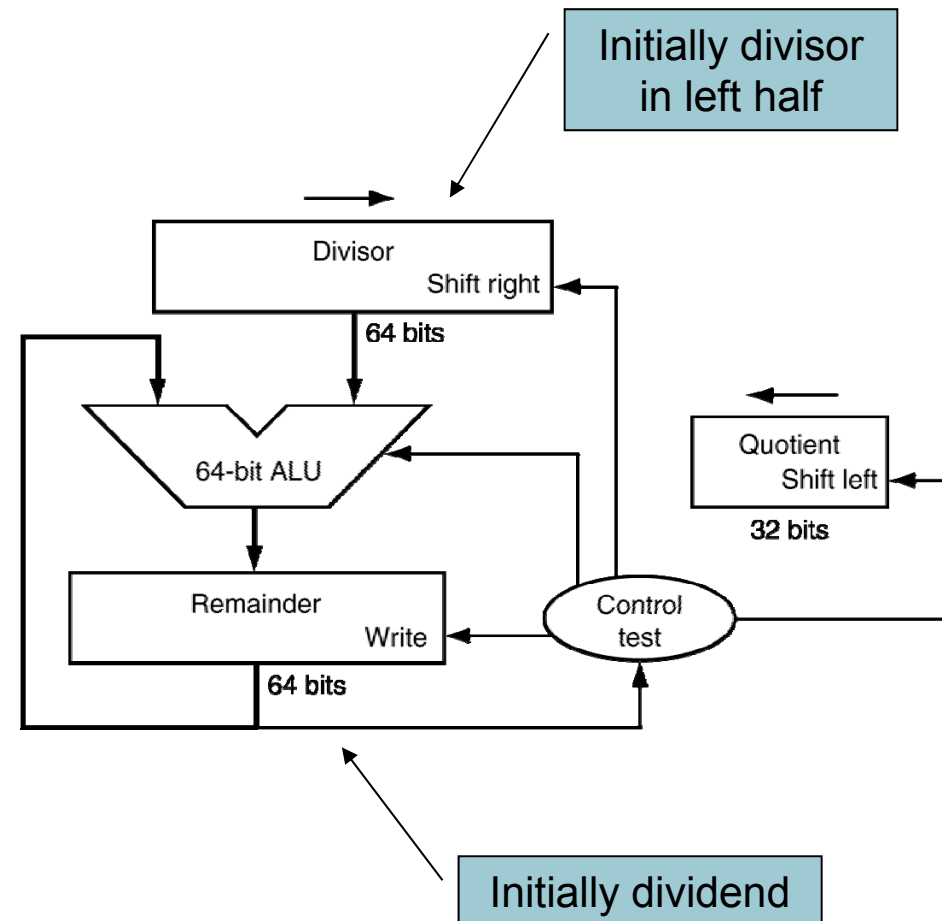
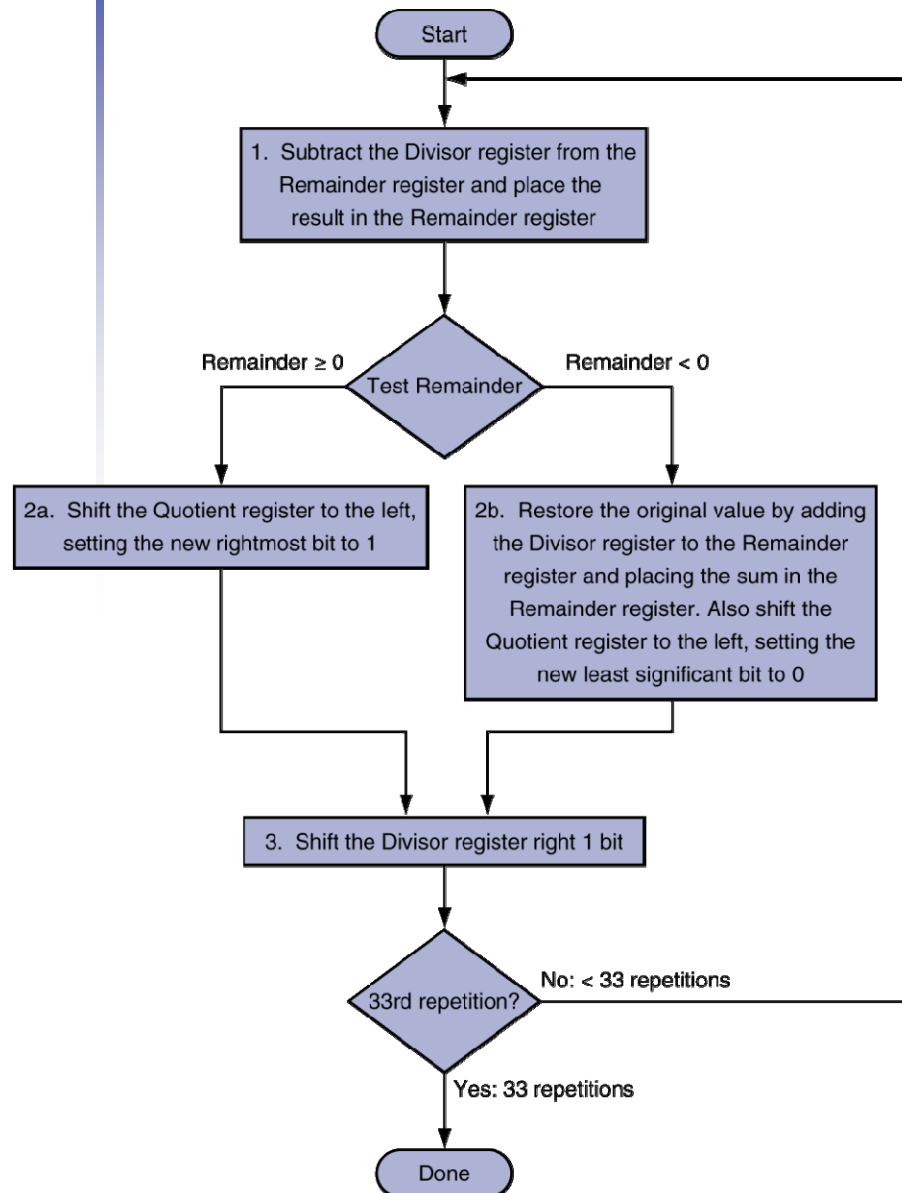


*n*-bit operands yield *n*-bit quotient and remainder

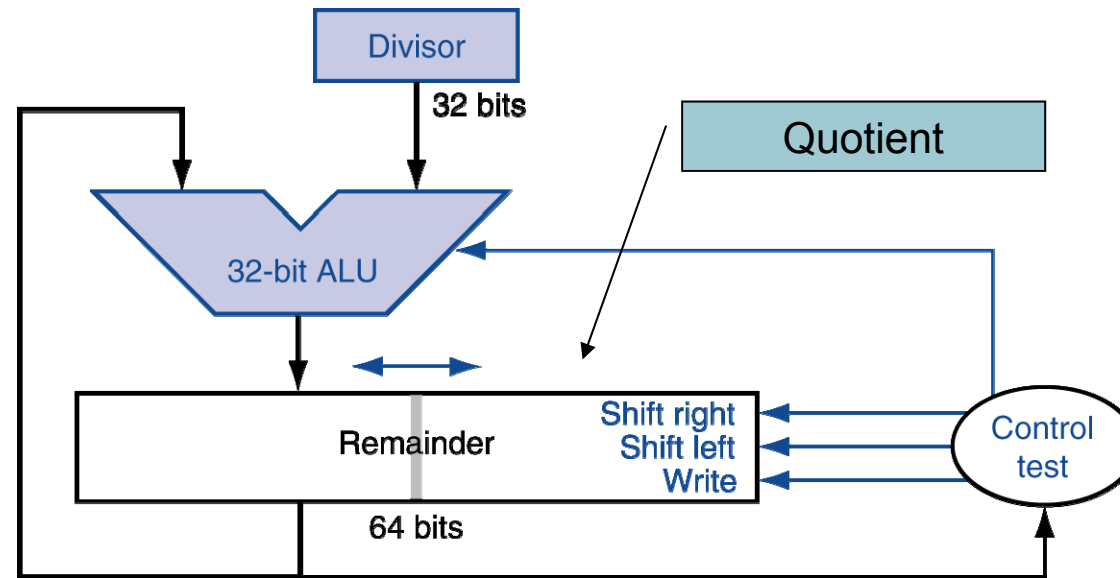
- Check for 0 divisor
- Long division approach
  - If divisor  $\leq$  dividend bits
    - 1 bit in quotient, subtract
  - Otherwise
    - 0 bit in quotient, bring down next dividend bit
- Restoring division
  - Do the subtract, and if remainder goes  $< 0$ , add divisor back
- Signed division
  - Divide using absolute values
  - Adjust sign of quotient and remainder as required



# Division Hardware



# Optimized Divider



- One cycle per partial-remainder subtraction
- Looks a lot like a multiplier!
  - Same hardware can be used for both

# Faster Division

- Can't use parallel hardware as in multiplier
  - Subtraction is conditional on sign of remainder
- Faster dividers (e.g. **SRT division**) generate multiple quotient bits per step
  - Guesses quotient bits using a table lookup based on upper bits of dividend, remainder
    - Subsequent steps correct wrong guesses
  - Still require multiple steps
- Other faster dividers exist
  - nonrestoring dividers, nonperforming dividers, ...

# MIPS Division

- Use HI/LO registers for result
  - **HI**: 32-bit remainder
  - **LO**: 32-bit quotient
- Instructions
  - `div rs, rt` / `divu rs, rt`
  - Use `mfhi` , `mflo` to access result
  - No overflow or divide-by-0 checking
    - Software must perform checks if required

# Floating Point

- Representation for non-integral numbers
  - Including very small and very large numbers
- Like scientific notation
  - $-2.34 \times 10^{56}$  ← normalized
  - $+0.002 \times 10^{-4}$  ← not normalized
  - $+987.02 \times 10^9$  ← not normalized
- In binary
  - $\pm 1.xxxxxxx_2 \times 2^{yyyy}$
- Types **float** and **double** in C

# Floating Point Standard

- Defined by IEEE 754 Standard
- Developed in response to divergence of representations
  - Portability issues for scientific code
- Now almost universally adopted
  - In every computer invented since 1980
- Two representations
  - **Single** precision (**32-bit**)
  - **Double** precision (**64-bit**)
    - reduces chance of **underflow** and **overflow**

# IEEE Floating-Point Format

single: 8 bits

double: 11 bits

single: 23 bits

double: 52 bits

S	Exponent	Fraction
---	----------	----------

$$x = (-1)^S \times (1 + \text{Fraction}) \times 2^{(\text{Exponent} - \text{Bias})}$$

- S: sign bit (0  $\Rightarrow$  non-negative, 1  $\Rightarrow$  negative)
- Normalize significand:  $1.0 \leq |\text{significand}| < 2.0$ 
  - Always has a leading pre-binary-point 1 bit, so no need to represent it explicitly (hidden bit)
  - Significand is Fraction with the “1.” restored
- Exponent: excess representation: **actual exponent + Bias**
  - Ensures exponent is unsigned
  - Single: Bias = 127; Double: Bias = 1203

# Single-Precision Range

- Exponents 00000000 and 11111111 reserved
- **Smallest value**
  - Exponent: 00000001  
 $\Rightarrow$  actual exponent =  $1 - 127 = -126$
  - Fraction: 000...00  $\Rightarrow$  significand = 1.0
  - $\pm 1.0 \times 2^{-126} \approx \pm 1.2 \times 10^{-38}$
- **Largest value**
  - exponent: 11111110  
 $\Rightarrow$  actual exponent =  $254 - 127 = +127$
  - Fraction: 111...11  $\Rightarrow$  significand  $\approx 2.0$
  - $\pm 2.0 \times 2^{+127} \approx \pm 3.4 \times 10^{+38}$



# Double-Precision Range

- Exponents 0000...00 and 1111...11 reserved
- **Smallest value**
  - Exponent: 000000000001  
⇒ actual exponent =  $1 - 1023 = -1022$
  - Fraction: 000...00 ⇒ significand = 1.0
  - $\pm 1.0 \times 2^{-1022} \approx \pm 2.2 \times 10^{-308}$
- **Largest value**
  - Exponent: 111111111110  
⇒ actual exponent =  $2046 - 1023 = +1023$
  - Fraction: 111...11 ⇒ significand  $\approx 2.0$
  - $\pm 2.0 \times 2^{+1023} \approx \pm 1.8 \times 10^{+308}$

# Floating-Point Precision

## ■ Relative precision

- all fraction bits are significant
- **Single:** approx  $2^{-23}$ 
  - Equivalent to  $23 \times \log_{10} 2 \approx 23 \times 0.3 \approx 6$  decimal digits of precision
- **Double:** approx  $2^{-52}$ 
  - Equivalent to  $52 \times \log_{10} 2 \approx 52 \times 0.3 \approx 16$  decimal digits of precision

Single precision		Double precision		Object represented
Exponent	Fraction	Exponent	Fraction	
0	0	0	0	0
0	Nonzero	0	Nonzero	$\pm$ denormalized number
1–254	Anything	1–2046	Anything	$\pm$ floating-point number
255	0	2047	0	$\pm$ infinity
255	Nonzero	2047	Nonzero	NaN (Not a Number)

# Floating-Point Example

- Represent  $-0.75$ 
  - $-0.75 = -0.11_2 = (-1)^1 \times 1.1_2 \times 2^{-1}$
  - $S = 1$
  - Fraction =  $1000\dots00_2$
  - Exponent =  $-1 + \text{Bias}$ 
    - Single:  $-1 + 127 = 126 = 01111110_2$
    - Double:  $-1 + 1023 = 1022 = 011111111110_2$
- Single:  $10111111101000\dots00$
- Double:  $101111111111101000\dots00$

# Floating-Point Example

- What number is represented by the single-precision float

11000000101000...00

- $S = 1$

- Fraction =  $01000...00_2$

- Exponent =  $10000001_2 = 129$

- $$\begin{aligned}x &= (-1)^1 \times (1 + .01_2) \times 2^{(129 - 127)} \\&= (-1) \times 1.25 \times 2^2 \\&= -5.0\end{aligned}$$

# Infinites and NaNs

- Exponent = 111...1, Fraction = 000...0
  - $\pm$ Infinity
  - Can be used in subsequent calculations, avoiding need for overflow check
- Exponent = 111...1, Fraction  $\neq$  000...0
  - Not-a-Number (NaN)
  - Indicates illegal or undefined result
    - e.g., 0.0 / 0.0
  - Can be used in subsequent calculations

# Floating-Point Addition

- Consider a 4-digit decimal example
  - $9.999 \times 10^1 + 1.610 \times 10^{-1}$
- 1. Align decimal points
  - Shift number with smaller exponent
  - $9.999 \times 10^1 + 0.016 \times 10^1$
- 2. Add significands
  - $9.999 \times 10^1 + 0.016 \times 10^1 = 10.015 \times 10^1$
- 3. Normalize result & check for over/underflow
  - $1.0015 \times 10^2$
- 4. Round and renormalize if necessary
  - $1.002 \times 10^2$

# Floating-Point Addition

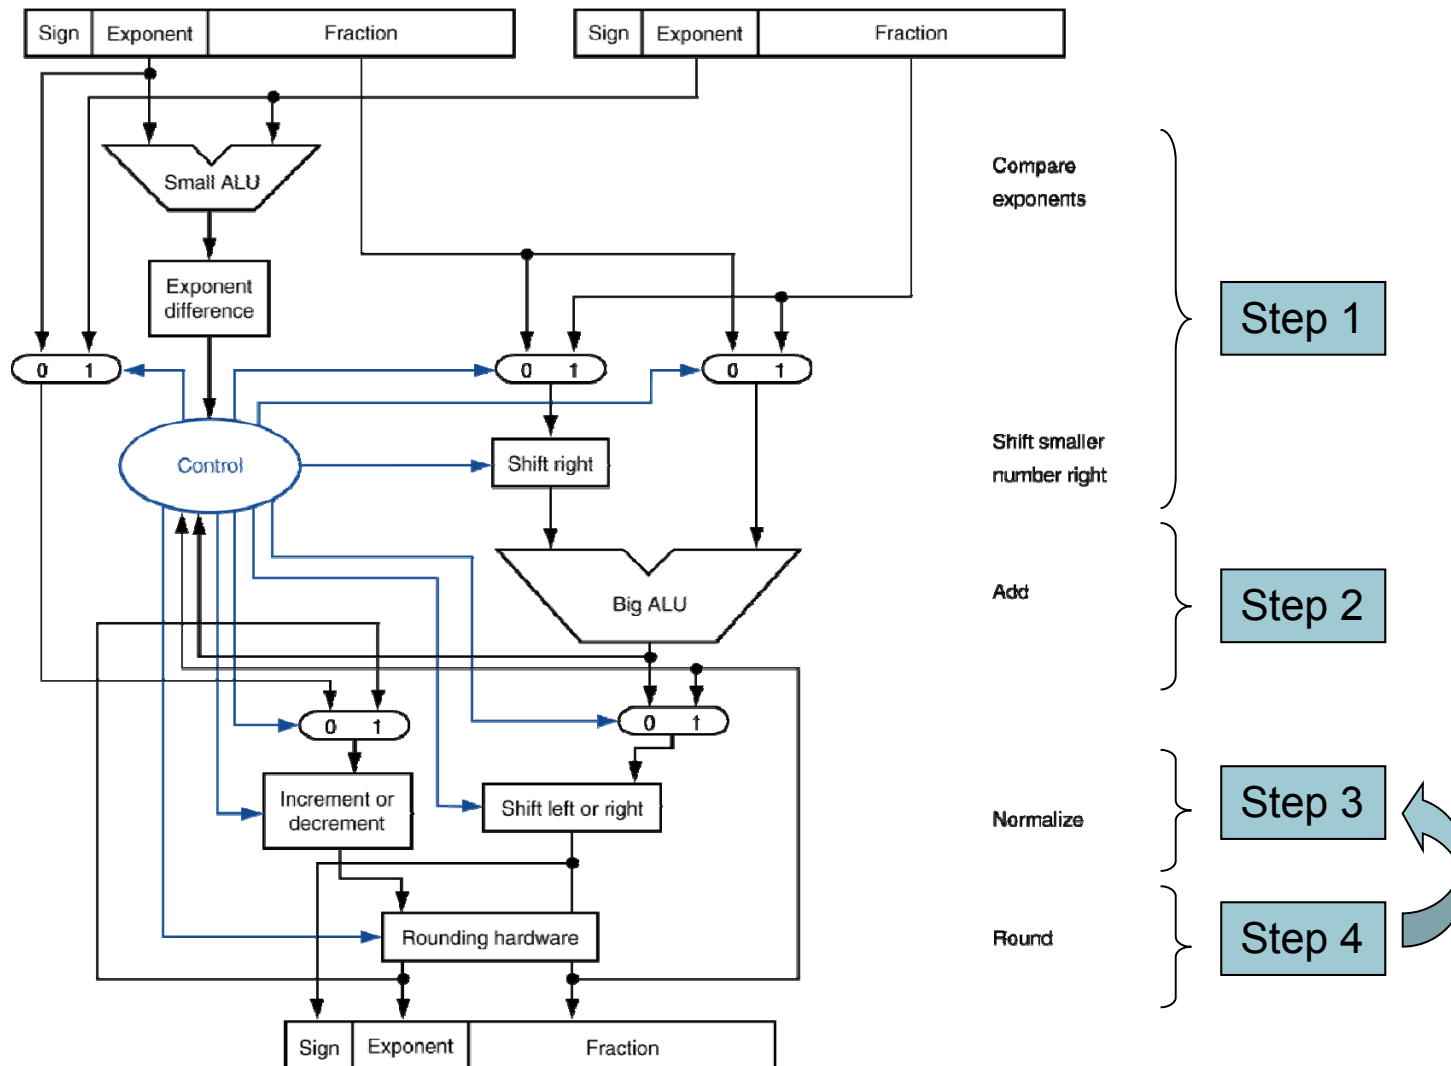
- Now consider a 4-digit binary example
  - $1.000_2 \times 2^{-1} + -1.110_2 \times 2^{-2}$  (0.5 + -0.4375)
- 1. Align binary points
  - Shift number with smaller exponent
  - $1.000_2 \times 2^{-1} + -0.111_2 \times 2^{-1}$
- 2. Add significands
  - $1.000_2 \times 2^{-1} + -0.111_2 \times 2^{-1} = 0.001_2 \times 2^{-1}$
- 3. Normalize result & check for over/underflow
  - $1.000_2 \times 2^{-4}$ , with no over/underflow
- 4. Round and renormalize if necessary
  - $1.000_2 \times 2^{-4}$  (no change) = 0.0625

# FP Adder Hardware

- Much more complex than integer adder
- Doing it in one clock cycle would take too long
  - Much longer than integer operations
  - Slower clock would penalize all instructions
- FP adder usually takes several cycles
  - Can be pipelined



# FP Adder Hardware



# Floating-Point Multiplication

- Consider a 4-digit decimal example
  - $1.110 \times 10^{10} \times 9.200 \times 10^{-5}$
- 1. Add exponents
  - For biased exponents, subtract bias from sum
  - New exponent =  $10 + -5 = 5$
- 2. Multiply significands
  - $1.110 \times 9.200 = 10.212 \Rightarrow 10.212 \times 10^5$
- 3. Normalize result & check for over/underflow
  - $1.0212 \times 10^6$
- 4. Round and renormalize if necessary
  - $1.021 \times 10^6$
- 5. Determine sign of result from signs of operands
  - $+1.021 \times 10^6$

# Floating-Point Multiplication

- Now consider a 4-digit binary example
  - $1.000_2 \times 2^{-1} \times -1.110_2 \times 2^{-2}$  ( $0.5 \times -0.4375$ )
- 1. Add exponents
  - Unbiased:  $-1 + -2 = -3$
  - Biased:  $(-1 + 127) + (-2 + 127) = -3 + 254 - 127 = -3 + 127$
- 2. Multiply significands
  - $1.000_2 \times 1.110_2 = 1.1102 \Rightarrow 1.110_2 \times 2^{-3}$
- 3. Normalize result & check for over/underflow
  - $1.110_2 \times 2^{-3}$  (no change) with no over/underflow
- 4. Round and renormalize if necessary
  - $1.110_2 \times 2^{-3}$  (no change)
- 5. Determine sign:  $+ve \times -ve \Rightarrow -ve$ 
  - $-1.110_2 \times 2^{-3} = -0.21875$

# FP Arithmetic Hardware

- FP multiplier is of similar complexity to FP adder
  - But uses a multiplier for significands instead of an adder
- FP arithmetic hardware usually does
  - Addition, subtraction, multiplication, division, reciprocal, square-root
  - $FP \leftrightarrow$  integer conversion
- Operations usually takes several cycles
  - Can be pipelined

# FP Instructions in MIPS

- FP hardware is coprocessor 1
  - Adjunct processor that extends the ISA
- Separate FP registers
  - 32 single-precision: `$f0, $f1, ... $f31`
  - Paired for double-precision: `$f0/$f1, $f2/$f3, ...`
    - Release 2 of MIPS ISA supports  $32 \times 64$ -bit FP reg's
- FP instructions operate only on FP registers
  - Programs generally don't do integer ops on FP data, or vice versa
  - More registers with minimal code-size impact
- FP load and store instructions
  - `lwc1, ldc1, swc1, sdc1`
    - e.g., `ldc1 $f8, 32($sp)`

# FP Instructions in MIPS

- Single-precision arithmetic
  - `add.s, sub.s, mul.s, div.s`
    - e.g., `add.s $f0, $f1, $f6`
- Double-precision arithmetic
  - `add.d, sub.d, mul.d, div.d`
    - e.g., `mul.d $f4, $f4, $f6`
- Single- and double-precision comparison
  - `c.xx.s, c.xx.d` (`xx` is `eq, lt, le, ...`)
  - Sets or clears FP condition-code bit
    - e.g. `c.lt.s $f3, $f4`
- Branch on FP condition code true or false
  - `bc1t, bc1f`
    - e.g., `bc1t TargetLabel`

# FP Example: °F to °C

- C code:

```
float f2c (float fahr) {  
    return ((5.0/9.0)*(fahr - 32.0));  
}
```

- fahr in **\$f12**, result in **\$f0**, literals in global memory space

- Compiled MIPS code:

```
f2c: lw1    $f16, const5($gp)  
     lw2    $f18, const9($gp)  
     div.s  $f16, $f16, $f18  
     lw1    $f18, const32($gp)  
     sub.s  $f18, $f12, $f18  
     mul.s  $f0, $f16, $f18  
     jr     $ra
```

# FP Example: Array Multiplication

- $X = X + Y \times Z$ 
  - All  $32 \times 32$  matrices, 64-bit double-precision elements

- C code:

```
void mm (double x[],  
         double y[], double z[]) {  
    int i, j, k;  
    for (i = 0; i != 32; i = i + 1)  
        for (j = 0; j != 32; j = j + 1)  
            for (k = 0; k != 32; k = k + 1)  
                x[i][j] = x[i][j]  
                    + y[i][k] * z[k][j];  
}
```

- Addresses of x, y, z in  $\$a0$ ,  $\$a1$ ,  $\$a2$ , and  
i, j, k in  $\$s0$ ,  $\$s1$ ,  $\$s2$



# FP Example: Array Multiplication

## ■ MIPS code:

	li	\$t1, 32	# \$t1 = 32 (row size/loop end)
	li	\$s0, 0	# i = 0; initialize 1st for loop
L1:	li	\$s1, 0	# j = 0; restart 2nd for loop
L2:	li	\$s2, 0	# k = 0; restart 3rd for loop
	sll	\$t2, \$s0, 5	# \$t2 = i * 32 (size of row of x)
	addu	\$t2, \$t2, \$s1	# \$t2 = i * size(row) + j
	sll	\$t2, \$t2, 3	# \$t2 = byte offset of [i][j]
	addu	\$t2, \$a0, \$t2	# \$t2 = byte address of x[i][j]
	l.d	\$f4, 0(\$t2)	# \$f4 = 8 bytes of x[i][j]
L3:	sll	\$t0, \$s2, 5	# \$t0 = k * 32 (size of row of z)
	addu	\$t0, \$t0, \$s1	# \$t0 = k * size(row) + j
	sll	\$t0, \$t0, 3	# \$t0 = byte offset of [k][j]
	addu	\$t0, \$a2, \$t0	# \$t0 = byte address of z[k][j]
	l.d	\$f16, 0(\$t0)	# \$f16 = 8 bytes of z[k][j]

...

# FP Example: Array Multiplication

...

sll	\$t0, \$s0, 5	# \$t0 = i * 32 (size of row of y)
addu	\$t0, \$t0, \$s2	# \$t0 = i * size(row) + k
sll	\$t0, \$t0, 3	# \$t0 = byte offset of [i][k]
addu	\$t0, \$a1, \$t0	# \$t0 = byte address of y[i][k]
l.d	\$f18, 0(\$t0)	# \$f18 = 8 bytes of y[i][k]
mul.d	\$f16, \$f18, \$f16	# \$f16 = y[i][k] * z[k][j]
add.d	\$f4, \$f4, \$f16	# f4 = x[i][j] + y[i][k] * z[k][j]
addiu	\$s2, \$s2, 1	# \$k = k + 1
bne	\$s2, \$t1, L3	# if (k != 32) go to L3
s.d	\$f4, 0(\$t2)	# x[i][j] = \$f4
addiu	\$s1, \$s1, 1	# \$j = j + 1
bne	\$s1, \$t1, L2	# if (j != 32) go to L2
addiu	\$s0, \$s0, 1	# \$i = i + 1
bne	\$s0, \$t1, L1	# if (i != 32) go to L1

# Accurate Arithmetic

- Infinite variety of real numbers between, say, 0 and 1
  - Only  $2^{53}$  can be represented by double precision FP
- IEEE Std 754 specifies additional rounding control
  - Extra bits of precision (**guard, round, sticky**)
  - **guard** and **round** bits are 2 extra bits kept on the right during intermediate additions
  - **sticky** bit used in rounding in addition to **guard** and **round** bits; is set whenever a 1 bit shifts right of the round bit

**F** = 1 . xxxxxxxxxxxxxxxxxxxxxxxxxxxx **G R S**

- Not all FP units implement all options
  - Most programming languages and FP libraries just use defaults
- Trade-off between hardware complexity, performance, and market requirements

# Interpretation of Data

## The BIG Picture

- Bits have no inherent meaning
  - Same bits can represent a variety of objects
  - Interpretation depends on the instructions applied
- Computer representations of numbers
  - Finite range and precision
  - Programmers, computer systems must **minimize gap** between computer arithmetic and real world arithmetic

# Parallelism and Associativity

- Parallel programs may interleave operations in unexpected orders
  - Integer addition is associative
  - Assumptions of associativity for FP numbers may fail!

		$(x+y)+z$	$x+(y+z)$
x	-1.50E+38	0.00E+00	-1.50E+38
y	1.50E+38		1.50E+38
z	1.0		
		1.00E+00	0.00E+00

- Floating point numbers are approximations of real numbers – not associative!
- Need to validate parallel programs under varying degrees of parallelism

# x86 FP Architecture

- Originally based on 8087 FP coprocessor
  - 8 × 80-bit extended-precision registers
  - Used as a push-down stack
  - Registers indexed from TOS: ST(0), ST(1), ...
- FP values are 32-bit or 64-bit in memory
  - Converted on load/store of memory operand
  - Integer operands can also be converted on load/store

# x86 FP Instructions

Data transfer	Arithmetic	Compare	Transcendental
FILD mem/ST(i) FISTP mem/ST(i) FLDPI FLD1 FLDZ	FADDP mem/ST(i) FSUBRP mem/ST(i) FMULP mem/ST(i) FDIVRP mem/ST(i) FSQRT FABS FRNDINT	FICOMP FIUCOMP FSTSW AX/mem	FPATAN F2XMI FCOS FPTAN FPREM FPSIN FYL2X

- No FP branch – FSTSW sends result of CMP to INT CPU
- Optional variations
  - **I**: integer operand
  - **P**: pop operand from stack
  - **R**: reverse operand order
  - But not all combinations allowed

# Streaming SIMD Extension 2 (SSE2)

- Adds 4 × 128-bit registers
  - Extended to 8 registers in AMD64/EM64T
- Can be used for multiple FP operands
  - 2 × 64-bit double precision
  - 4 × 32-bit single precision
  - Instructions operate on them simultaneously
    - Single-Instruction Multiple-Data



# Fallacy: Right Shift and Division

- Left shift by  $i$  places multiplies an integer by  $2^i$
- Right shift divides by  $2^i$ ?
  - Only for unsigned integers!
- For signed integers
  - Logical right shift is clearly erroneous
    - e.g.,  $-5 / 4$
    - $11111011_2 \ggg 2 = 00111110_2 = +62$
  - Arithmetic right shift - replicate the sign bit
    - $11111011_2 \gg 2 = 11111110_2 = -2$
    - Result is -2 instead of -1; close, but no cigar

# Who Cares About FP Accuracy?

- Important for scientific code
  - But for everyday consumer use?
    - “My bank balance is out by 0.0002¢!” ☹
- The Intel Pentium FDIV bug (~1994)
  - Bug in LUT used to guess multiple quotient bits per step; wrong values in some LUT locations
  - **Cost Intel \$300+ million**
  - The market expects accuracy
  - See Colwell, *The Pentium Chronicles*



# Summary: MIPS Instruction Set

MIPS core instructions	Name	Format	MIPS arithmetic core	Name	Format
add	add	R	multiply	mult	R
add immediate	addi	I	multiply unsigned	multu	R
add unsigned	addu	R	divide	div	R
add immediate unsigned	addiu	I	divide unsigned	divu	R
subtract	sub	R	move from Hi	mfhi	R
subtract unsigned	subu	R	move from Lo	mflo	R
AND	AND	R	move from system control (EPC)	mfc0	R
AND immediate	ANDi	I	floating-point add single	add.s	R
OR	OR	R	floating-point add double	add.d	R
OR immediate	ORi	I	floating-point subtract single	sub.s	R
NOR	NOR	R	floating-point subtract double	sub.d	R
shift left logical	sll	R	floating-point multiply single	mul.s	R
shift right logical	srl	R	floating-point multiply double	mul.d	R
load upper immediate	lui	I	floating-point divide single	div.s	R
load word	lw	I	floating-point divide double	div.d	R
store word	sw	I	load word to floating-point single	lwc1	I
load halfword unsigned	lhu	I	store word to floating-point single	swc1	I
store halfword	sh	I	load word to floating-point double	ldc1	I
load byte unsigned	lbu	I	store word to floating-point double	sdcl	I
store byte	sb	I	branch on floating-point true	bclt	I
load linked ( <i>atomic update</i> )	ll	I	branch on floating-point false	bclf	I
store cond. ( <i>atomic update</i> )	sc	I	floating-point compare single	c.x.s	R
branch on equal	beq	I	(x = eq, neq, lt, le, gt, ge)		
branch on not equal	bne	I	floating-point compare double	c.x.d	R
jump	j	J	(x = eq, neq, lt, le, gt, ge)		
jump and link	jal	J			
jump register	jr	R			
set less than	slt	R			
set less than immediate	slti	I			
set less than unsigned	sltu	R			
set less than immediate unsigned	sltiu	I			

# Summary: MIPS Instruction Set

Remaining MIPS-32	Name	Format	Pseudo MIPS	Name	Format
exclusive or ( $rs \oplus rt$ )	xor	R	absolute value	abs	rd,rs
exclusive or immediate	xori	I	negate ( <i>signed or unsigned</i> )	neg <sub>s</sub>	rd,rs
shift right arithmetic	sra	R	rotate left	rol	rd,rs,rt
shift left logical variable	sllv	R	rotate right	ror	rd,rs,rt
shift right logical variable	srlv	R	multiply and don't check oflw ( <i>signed or uns.</i> )	mul <sub>s</sub>	rd,rs,rt
shift right arithmetic variable	srav	R	multiply and check oflw ( <i>signed or uns.</i> )	mulos <sub>s</sub>	rd,rs,rt
move to Hi	mthi	R	divide and check overflow	div	rd,rs,rt
move to Lo	mtlo	R	divide and don't check overflow	divu	rd,rs,rt
load halfword	lh	I	remainder ( <i>signed or unsigned</i> )	rem <sub>s</sub>	rd,rs,rt
load byte	lb	I	load immediate	li	rd,imm
load word left ( <i>unaligned</i> )	lwl	I	load address	la	rd,addr
load word right ( <i>unaligned</i> )	lwr	I	load double	ld	rd,addr
store word left ( <i>unaligned</i> )	swl	I	store double	sd	rd,addr
store word right ( <i>unaligned</i> )	swr	I	unaligned load word	ulw	rd,addr
load linked ( <i>atomic update</i> )	ll	I	unaligned store word	usw	rd,addr
store cond. ( <i>atomic update</i> )	sc	I	unaligned load halfword ( <i>signed or uns.</i> )	ulh <sub>s</sub>	rd,addr
move if zero	movz	R	unaligned store halfword	ush	rd,addr
move if not zero	movn	R	branch	b	Label
multiply and add ( <i>S or uns.</i> )	madd <sub>s</sub>	R	branch on equal zero	beqz	rs,L
multiply and subtract ( <i>S or uns.</i> )	msub <sub>s</sub>	I	branch on compare ( <i>signed or unsigned</i> )	bx <sub>s</sub>	rs,rt,L
branch on $\geq$ zero and link	bgezal	I	( $x = lt, le, gt, ge$ )		
branch on $<$ zero and link	bltzal	I	set equal	seq	rd,rs,rt
jump and link register	jalr	R	set not equal	sne	rd,rs,rt
branch compare to zero	bxz	I	set on compare ( <i>signed or unsigned</i> )	sx <sub>s</sub>	rd,rs,rt
branch compare to zero likely	bxzl	I	( $x = lt, le, gt, ge$ )		
( $x = lt, le, gt, ge$ )			load to floating point ( <i>s or d</i> )	l. <sub>f</sub>	rd,addr
branch compare reg likely	bxll	I	store from floating point ( <i>s or d</i> )	s. <sub>f</sub>	rd,addr
trap if compare reg	tx	R			
trap if compare immediate	txi	I			
( $x = eq, neq, lt, le, gt, ge$ )					
return from exception	rfe	R			
system call	syscall	I			
break ( <i>cause exception</i> )	break	I			
move from FP to integer	mfc1	R			
move to FP from integer	mtc1	R			
FP move ( <i>s or d</i> )	mov. <sub>f</sub>	R			
FP move if zero ( <i>s or d</i> )	movz. <sub>f</sub>	R			
FP move if not zero ( <i>s or d</i> )	movn. <sub>f</sub>	R			
FP square root ( <i>s or d</i> )	sqr. <sub>f</sub>	R			
FP absolute value ( <i>s or d</i> )	abs. <sub>f</sub>	R			
FP negate ( <i>s or d</i> )	neg. <sub>f</sub>	R			
FP convert ( <i>w, s, or d</i> )	cvt. <sub>ff</sub>	R			
FP compare un ( <i>s or d</i> )	c.xn. <sub>f</sub>	R			

# Frequency of Common MIPS Instructions

- Only included those with >3% (table 1) and >1% (table 2)

<b>MIPS core</b>	<b>SPECint</b>	<b>SPECfp</b>
addu	5.2%	3.5%
addiu	9.0%	7.2%
or	4.0%	1.2%
sll	4.4%	1.9%
lui	3.3%	0.5%
lw	18.6%	5.8%
sw	7.6%	2.0%
lbu	3.7%	0.1%
beq	8.6%	2.2%
bne	8.4%	1.4%
slt	9.9%	2.3%
slti	3.1%	0.3%
sltu	3.4%	0.8%

<b>Arith core + MIPS-32</b>	<b>SPECint</b>	<b>SPECfp</b>
add.d	0.0%	10.6%
sub.d	0.0%	4.9%
mul.d	0.0%	15.0%
add.s	0.0%	1.5%
sub.s	0.0%	1.8%
mul.s	0.0%	2.4%
l.d	0.0%	17.5%
s.d	0.0%	4.9%
l.s	0.0%	4.2%
s.s	0.0%	1.1%
lhu	1.3%	0.0%

# Concluding Remarks

- ISAs support arithmetic
  - Signed and unsigned integers
  - Floating-point approximation for reals
- Bounded range and precision
  - Operations can overflow and underflow
- MIPS ISA
  - MIPS core and arithmetic core instructions: 54 most frequently used
    - 100% of SPECINT, 97% of SPECFP
  - Other instructions: less frequent